

First, connect to HBase:

Need to install MySQL connector before starting with Apache Sqoop to use Sqoop with MySQL tables

Copying it to the Sqoop library to complete the installation.

```

mysql-connector-java-8.0.25/src/test/java/testsuite/simple/SimpleTransformer.java
mysql-connector-java-8.0.25/src/test/java/testsuite/simple/SplitDBdotNameTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/simple/StatementsTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/simple/StringUtilsTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/simple/TestLifecycleInterceptor.java
mysql-connector-java-8.0.25/src/test/java/testsuite/simple/TransactionTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/simple/TraversalTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/simple/UpdatabilityTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/simple/UtilsTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/simple/XATest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/AsyncTests.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/BaseXDevAPITestCase.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/TestXDevAPIRequirements.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/AsyncQueryTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/BaseCollectionTestCase.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/BaseTableTestCase.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/BindTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/CollectionAddTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/CollectionFindTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/CollectionModifyTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/CollectionRemoveTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/CollectionTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/CompressionTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/DevApiBaseTestCase.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/IPv6SupportTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/MetadataTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/ResultTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/SchemaTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/SecureSessionTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/SessionFailoverTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/SessionTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/TableDeleteTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/TableInsertTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/TableSelectTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/TableTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/TableUpdateTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/TransactionTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/devapi/package-info.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/InternalXBaseTestCase.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/MysqlxSessionTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/XProtocolAsyncTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/XProtocolAuthTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/XProtocolTest.java
mysql-connector-java-8.0.25/src/test/java/testsuite/x/internal/package-info.java
[hadoop@ip-172-31-30-111 ~]$ cd mysql-connector-java-8.0.25/
[hadoop@ip-172-31-30-111 mysql-connector-java-8.0.25]$ sudo cp mysql-connector-java-8.0.25.jar /usr/lib/sqoop/lib/
[hadoop@ip-172-31-30-111 mysql-connector-java-8.0.25]$

```

Use sqoop command to import data from MySQL to HBase:

```

at org.apache.sqoop.tool.ImportTool.run(ImportTool.java:625)
at org.apache.sqoop.Sqoop.run(Sqoop.java:147)
at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:76)
at org.apache.sqoop.Sqoop.run(Sqoop.java:183)
at org.apache.sqoop.Sqoop.runTool(Sqoop.java:234)
at org.apache.sqoop.Sqoop.runTool(Sqoop.java:243)
at org.apache.sqoop.Sqoop.main(Sqoop.java:252)
[root@ip-172-31-30-111 ~]# sqoop import --connect jdbc:mysql://database-2.cx3ljnqiod.us-east-1.rds.amazonaws.com:3306/amr --table mysql_taxi --columns "VendorID,tpep_pickup_datetime,tpep_dropoff_datetime,passenger_count,trip_distance,RatecodeID,store_and_fwd_flag,FULocationID,DOLocationID,payment_type,fare_amount,extra_tax,tax,tip amount,tolls amount,improvement_surcharge,total amount,congestion_surcharge,airport_fee,Serial_num" --username admin --password 12345678 --hbase-table hbase_taxi --column-family TaxiData --hbase-row-key Serial_num --connection-manager org.apache.sqoop.manager.GenericJdbcManager --driver com.mysql.jdbc.Driver --connection-param file connection.params -m 1
Warning: /usr/lib/sqoop/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
23/01/15 09:49:22 INFO sqoop.sqoop: Running Sqoop version: 1.4.7
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/lib/hadoop/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/share/aws/redshift/jdbc/redshift-jdbc42-1.2.37.1061.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
23/01/15 09:49:22 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
23/01/15 09:49:22 INFO manager.SqlManager: Using default fetchSize of 1000
23/01/15 09:49:22 INFO tool.CodeGenTool: Beginning code generation
Loading class 'com.mysql.jdbc.Driver'. This is deprecated. The new driver class is 'com.mysql.cj.jdbc.Driver'. The driver is automatically registered via the SPI and manual loading of the driver class is generally unnecessary.
23/01/15 09:49:22 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM mysql_taxi AS t WHERE 1=0
23/01/15 09:49:22 INFO orm.CompilationManager: HADOOP MAPRED HOME is /usr/lib/hadoop-mapreduce
Note: /tmp/sqoop-root/compile/3f0524a9675e5c3139c4d46ff4c4901f/mysql_taxi.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
23/01/15 09:49:25 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-root/compile/3f0524a9675e5c3139c4d46ff4c4901f/mysql_taxi.jar
23/01/15 09:49:25 INFO mapreduce.ImportJobBase: Beginning import of mysql_taxi
23/01/15 09:49:26 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
23/01/15 09:49:26 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
23/01/15 09:49:28 WARN mapreduce.TableMapReduceUtil: The addDependencyJars(Configuration, Class...) method has been deprecated since it is easy to use incorrectly. Most users should rely on addDependencyJars(Job) instead. See HBASE-8386 for more details.
23/01/15 09:49:28 INFO client.RMProxy: Connecting to ResourceManager at ip-172-31-30-111.ec2.internal/172.31.30.111:8032
23/01/15 09:49:36 INFO mapreduce.JobSubmitter: Using read committed transaction isolation
23/01/15 09:49:36 INFO mapreduce.JobSubmitter: number of splits:1
23/01/15 09:49:36 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1673772473890_0001
23/01/15 09:49:37 INFO impl.YarnClientImpl: Submitted application application_1673772473890_0001
23/01/15 09:49:37 INFO mapreduce.Job: The url to track the job: http://ip-172-31-30-111.ec2.internal:20888/proxy/application_1673772473890_0001/
23/01/15 09:49:37 INFO mapreduce.Job: Running job: job_1673772473890_0001
23/01/15 09:49:48 INFO mapreduce.Job: Job job_1673772473890_0001 running in uber mode : false
23/01/15 09:49:48 INFO mapreduce.Job: map 0% reduce 0%

```

Status is in running state:

aws

Services

Search

[Alt+S]

N. Virginia

voclabs/user2295247=parshavshivnani@gmail.com @ 1882-4764-5654

Amazon EMR

EMR Studio

EMR Serverless New

EMR on EC2

Clusters

Notebooks

Git repositories

Security configurations

Block public access

VPC subnets

Events

EMR on EKS

Virtual clusters

Help

What's new

Amazon EMR has launched a new console experience. [Learn more](#) or [switch to the new console](#)

The following table lists web interfaces you can view on the task nodes:

Application	User interface URL
HDFS Data Node	<a href="http://ec2-000-000-000-000.compute-1.amazonaws.com:50075/">http://ec2-000-000-000-000.compute-1.amazonaws.com:50075/</a>
Node Manager	<a href="http://ec2-000-000-000-000.compute-1.amazonaws.com:8042/">http://ec2-000-000-000-000.compute-1.amazonaws.com:8042/</a>

High-level application history

Amazon EMR collects information from YARN applications on your cluster and keeps a summary of historical information for seven days after applications have completed. [Learn more](#)

Effective January 23, 2023, the 7-day, high-level application history will be discontinued. Instead, connect to the Persistent Application user interfaces, which provide greater access to off-cluster application UIs and application logs for up to 30 days after the application ends. [Learn more](#)

YARN applications (1)

Filter: All applications

Filter applications ...

1 applications (all loaded)

Application ID	Type	Action	Status	Start time (UTC+5:30)	Duration	Finish time
application_1673772473890_0001	MapReduce	mysql_Taxi.jar	Running	2023-01-15 15:19 (UTC+5:30)	1.8 min	

Finally, the import gets completed. Status is showing succeeded.

Amazon EMR

EMR Studio

EMR Serverless New

EMR on EC2

Clusters

Notebooks

Git repositories

Security configurations

Block public access

VPC subnets

Events

EMR on EKS

Virtual clusters

Help

What's new

Amazon EMR has launched a new console experience. [Learn more](#) or [switch to the new console](#)

status

Resource Manager <http://ec2-52-87-250-73.compute-1.amazonaws.com:8088/> SSH tunnel not enabled

The following table lists web interfaces you can view on the task nodes:

Application	User interface URL
HDFS Data Node	<a href="http://ec2-000-000-000-000.compute-1.amazonaws.com:50075/">http://ec2-000-000-000-000.compute-1.amazonaws.com:50075/</a>
Node Manager	<a href="http://ec2-000-000-000-000.compute-1.amazonaws.com:8042/">http://ec2-000-000-000-000.compute-1.amazonaws.com:8042/</a>

High-level application history

Amazon EMR collects information from YARN applications on your cluster and keeps a summary of historical information for seven days after applications have completed. [Learn more](#)

Effective January 23, 2023, the 7-day, high-level application history will be discontinued. Instead, connect to the Persistent Application user interfaces, which provide greater access to off-cluster application UIs and application logs for up to 30 days after the application ends. [Learn more](#)

YARN applications (1)

Filter: All applications

Filter applications ...

1 applications (all loaded)

Application ID	Type	Action	Status	Start time (UTC+5:30)	Duration	Finish time
application_1673772473890_0001	MapReduce	mysql_Taxi.jar	Succeeded	2023-01-15 15:19 (UTC+5:30)	25 min	2023-01-15 15:44 (UTC+5:30)